

COMP3221: Distributed Systems

Architectures & Processes

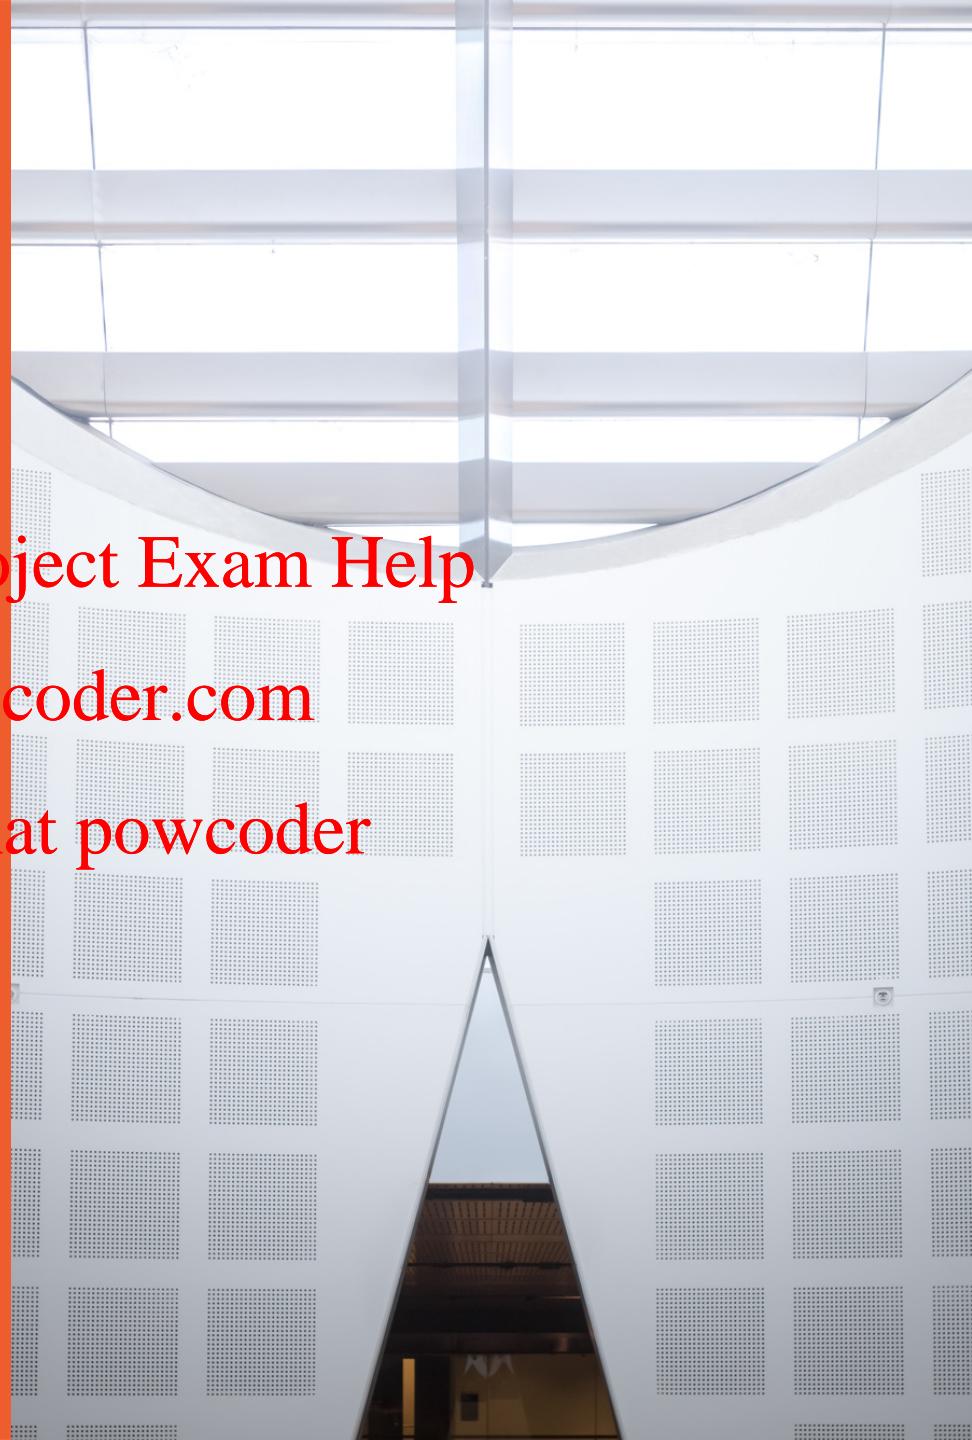
Assignment Project Exam Help

Add WeChat powcoder

Dr Nguyen Tran
School of Computer Science



THE UNIVERSITY OF
SYDNEY



Previously...

- Basic definition of a distributed system
 - “*A collection of independent computers that appears to its users as a single coherent system.*”

Assignment Project Exam Help

- Real-world examples for distributed systems

<https://powcoder.com>

- Transparency helps the users observe a single coherent system

Add WeChat powcoder

Transparency	Description
Access	Hide differences in data representation and how a resource is accessed
Location	Hide where a resource is located
Migration	Hide that a resource may move to another location
Relocation	Hide that a resource may be moved to another location while in use
Replication	Hide that a resource is replicated
Concurrency	Hide that a resource may be shared by several competitive users
Failure	Hide the failure and recovery of a resource

Outline - Architectures

- Software Architectures

- System Architectures

- Centralized architectures

- The Client-Server Model
 - The Layered Organization

- Decentralized architectures

- The Peer-to-Peer Organization

- Hybrid architectures

- Edge server systems
 - Collaborative distributed systems

Software Architectures

Architectures & Processes

Week 2, COMP3221

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



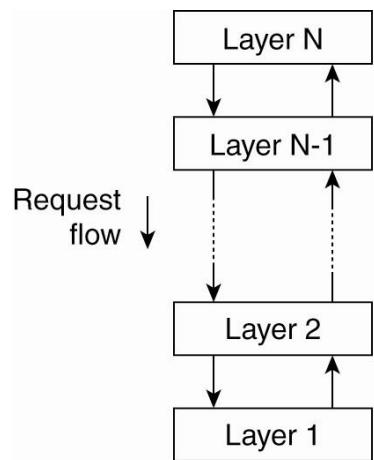
THE UNIVERSITY OF
SYDNEY

Software Architectures

Logical organization of components in distributed systems

- **Component:** A modular unit with well-defined interfaces that is replaceable
 - **Connector:** a mechanism that mediates communication, coordination among components.

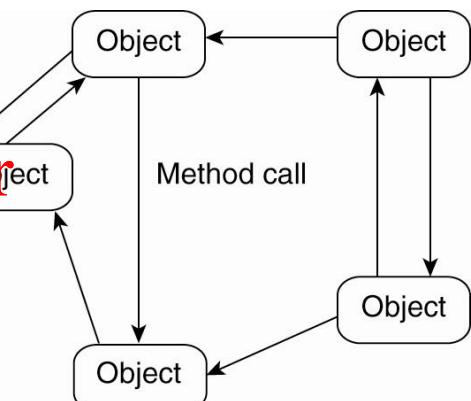
Layer-based Architectures vs. Object-based Architectures



(a)



Add WeChat powcoder



(b)

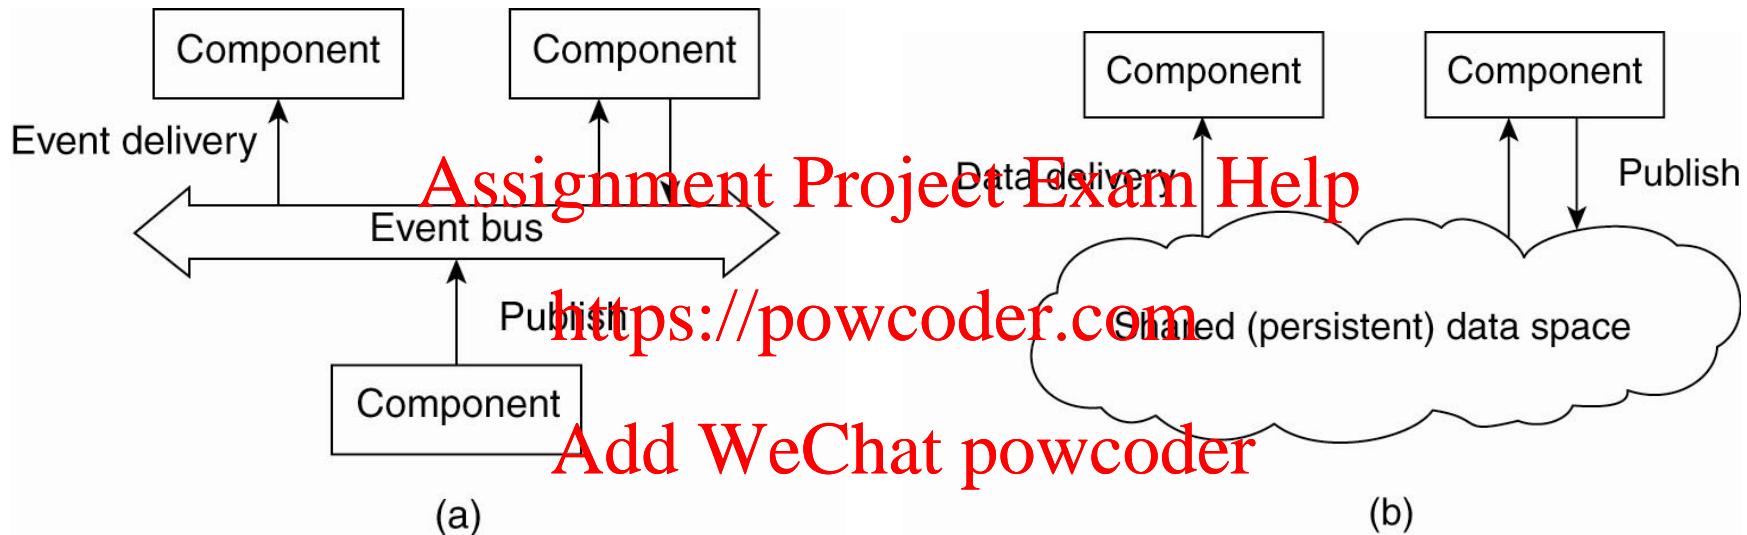
- In (a), requests (resp. responses) go downward (resp. upward)
 - In (b), objects communicate through Remote Procedure Calls (RPCs)

Software Architectures

Event-based Architectures

vs.

Data-centered Architectures



- (a) Event-based architecture: communication through events, that **optionally carry data** (subscribers get their desired events delivered)
- (b) Data-centered architecture: through a shared repository, that **contains data** (e.g., files in a distributed file system, web-based distributed system, Twitter)

Centralized Architectures

Assignment Project Exam Help

Architectures & Processes

Week 2, COMP3221

<https://powcoder.com>

Add WeChat powcoder

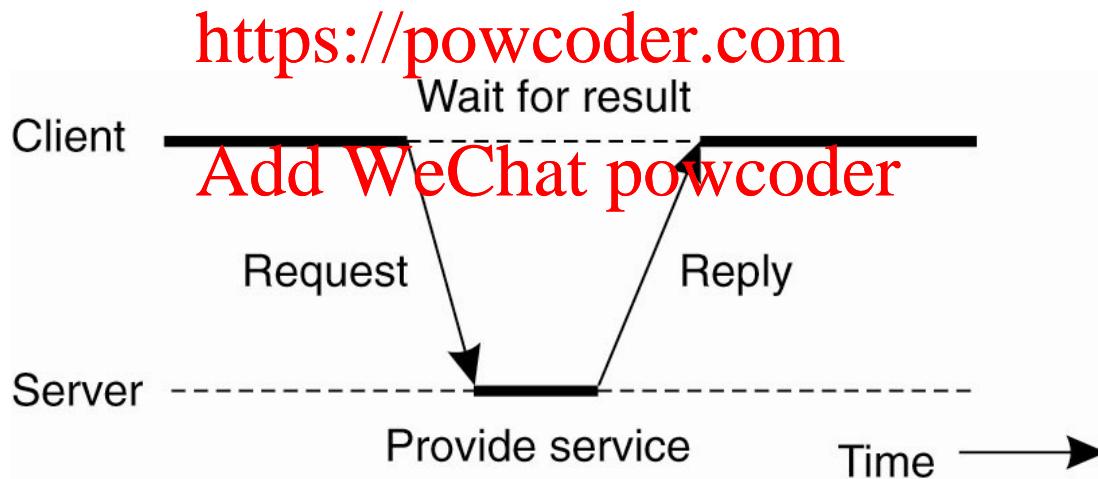


THE UNIVERSITY OF
SYDNEY

Centralized Architectures

The basic client-server model

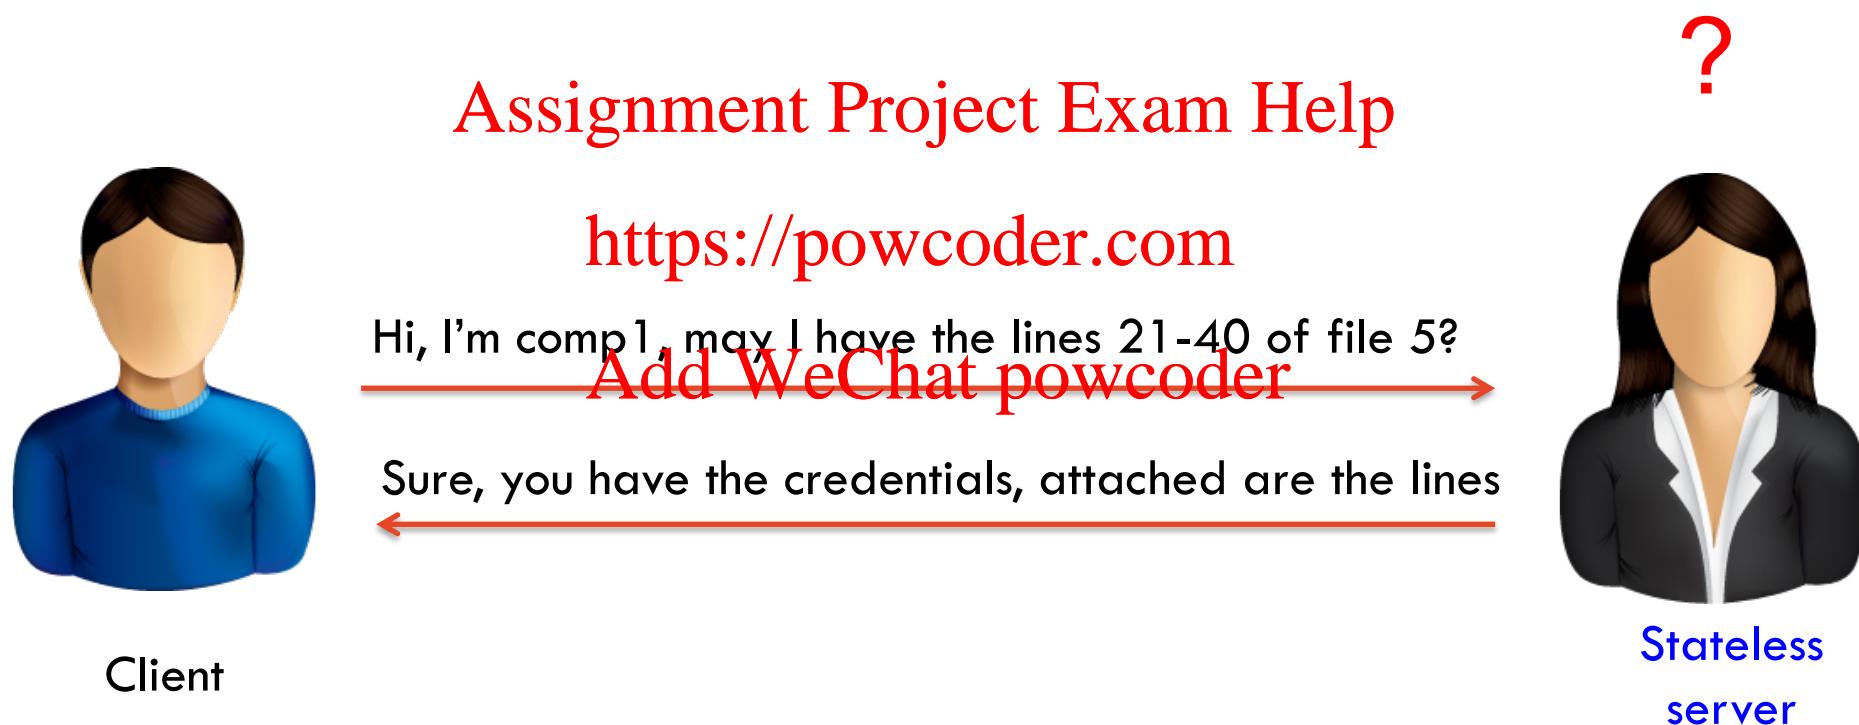
- The client/server requests/provides the services
- The client and the server can be hosted on different machines.



- The communication follows a request-reply model.
- Examples: ??

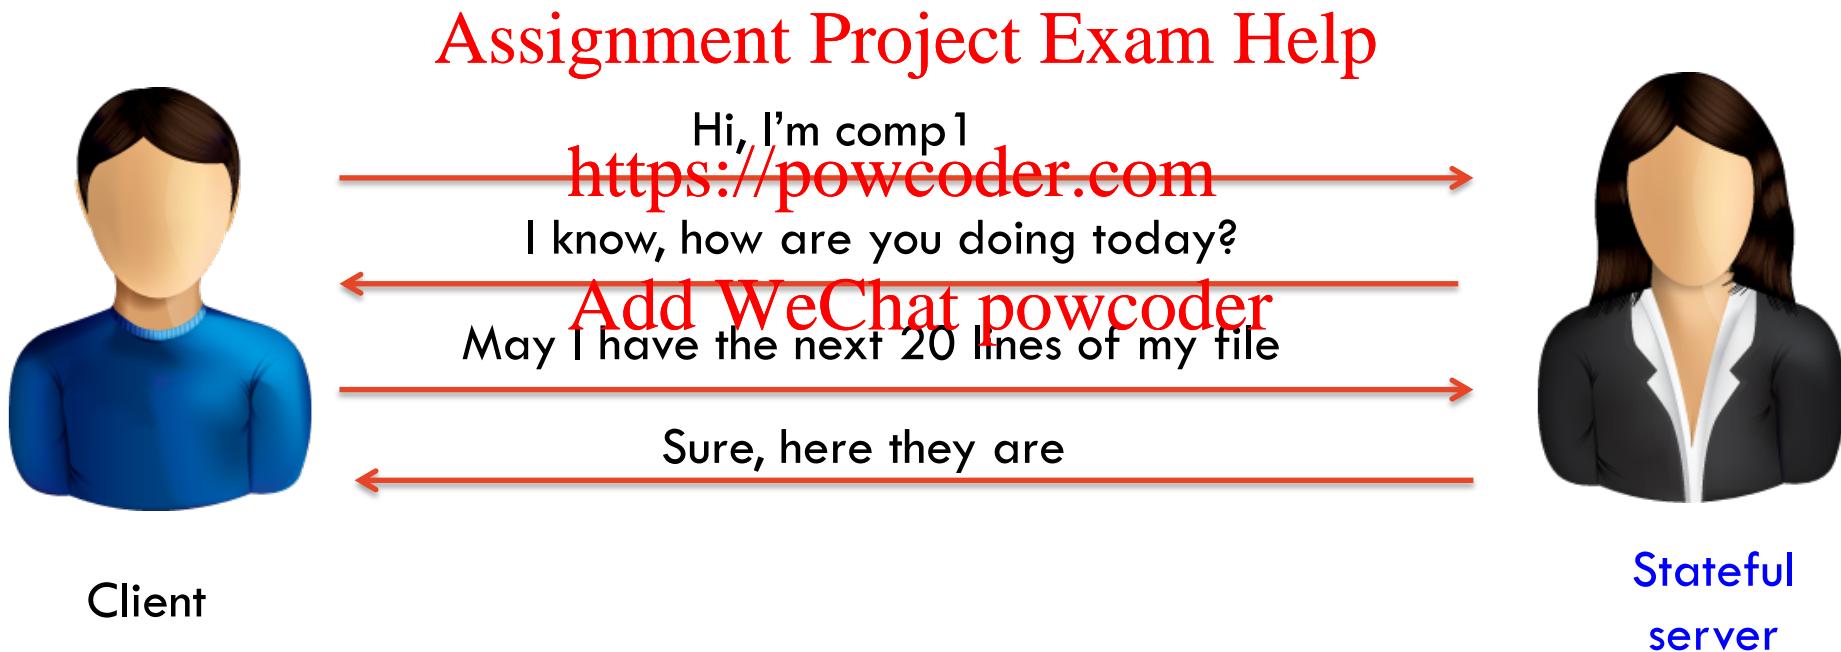
The Client-Server Model

Stateless vs Stateful server



The Client-Server Model

Stateless vs Stateful server



The Client-Server Model

Stateless vs Stateful server

- Stateless server: does not record the state of its clients
- Stateful server: maintains persistent information about its clients
(client->file)

<https://powcoder.com>

	Stateless server	Stateful server
State	No info kept	Persistent info
Request	Self-contained	Can be split, generally faster
Upon failure	No recovery needed	State recovery needed (explicit deletion)
Example	Network file system (NFSv3)	Andrew file system (AFS)

The Layered Organization

Architectures & Processes
Assignment Project Exam Help
Week 2, COMP3221
<https://powcoder.com>

Add WeChat powcoder

Application layering

Traditional three-layered view:

1. The user interface layer

- contains the feature to control the application

Assignment Project Exam Help

2. The processing layer

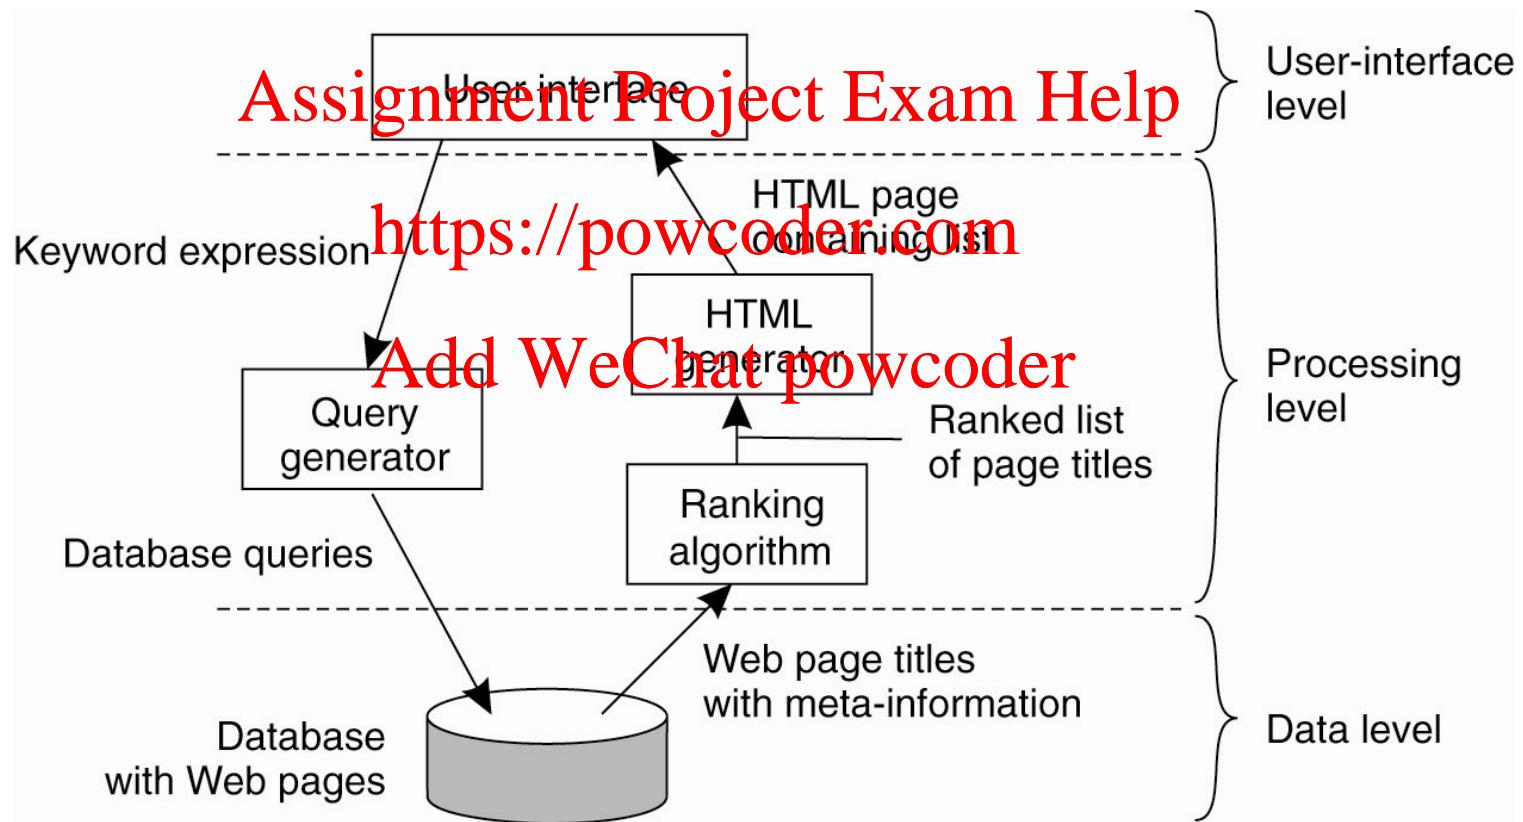
- contains the function of the application

3. The data layer

- contains the data of the application

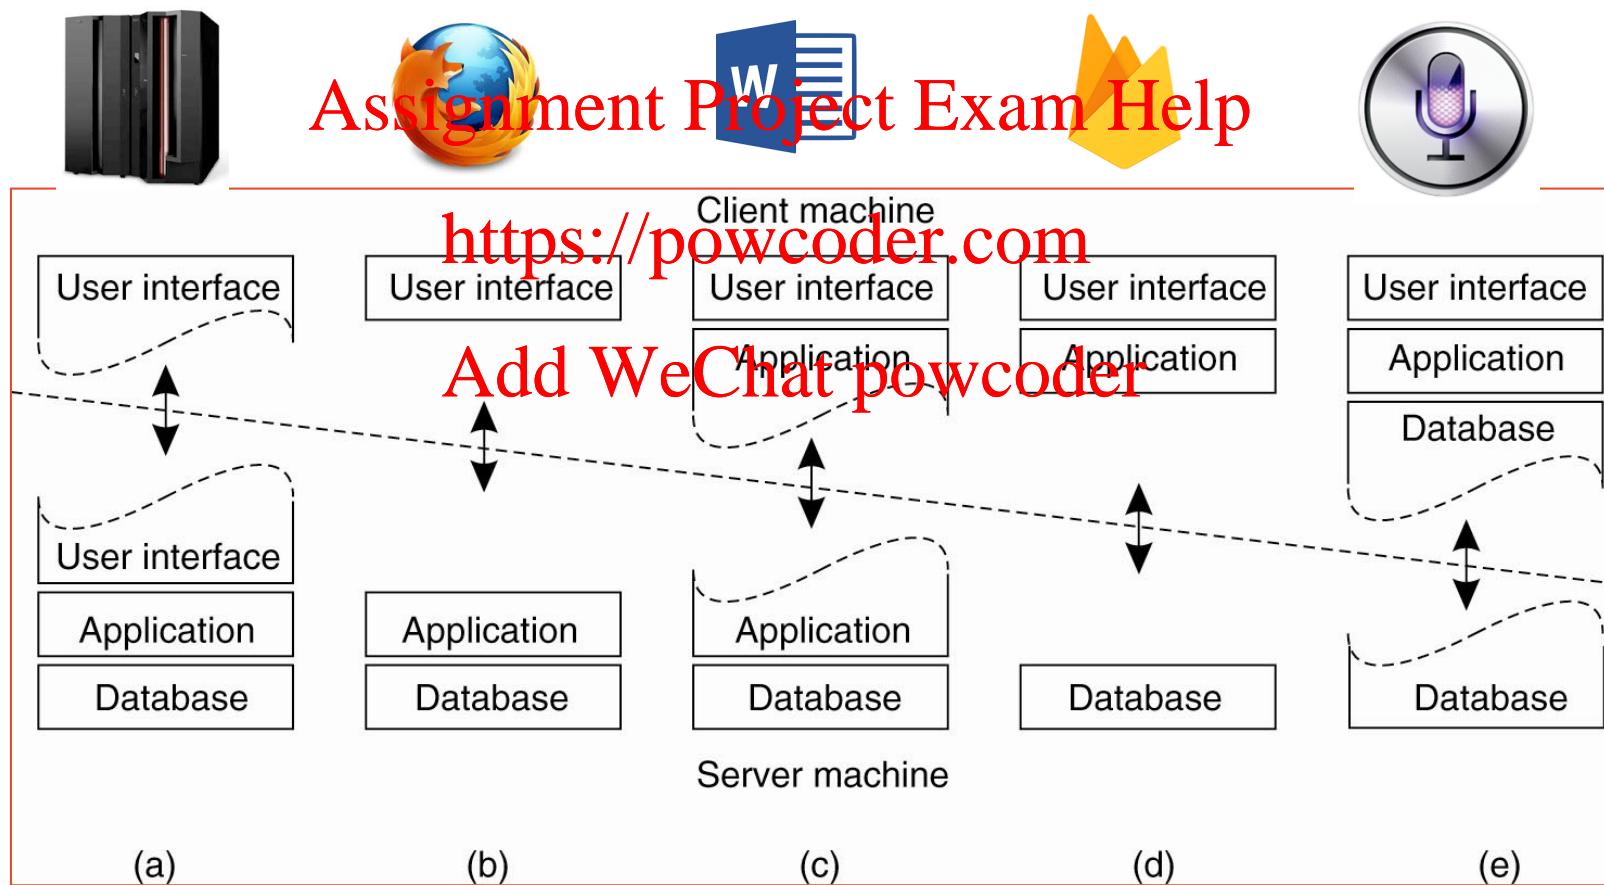
Application layering (cont'd)

Example: a search engine request spanning the traditional three layers



Multi-tiered architectures

Physical two-tiered architecture



Multi-tiered architectures (cont'd)

A single machine can act both as a client and a server

Example: Cloud computing

- *Cloud computing:* the delivery of computation or storage as a service to end-users.
- The servers handle most of the computation upon request and sends back the results to the client
 - One server asks data to another server
 - Another does the computation
 - ...

Assignment Project Exam Help

<https://powcoder.com>

amazon.com®



Decentralized Architectures

Assignment Project Exam Help

Architectures & Processes

Week 2, COMP3221

<https://powcoder.com>

Add WeChat powcoder

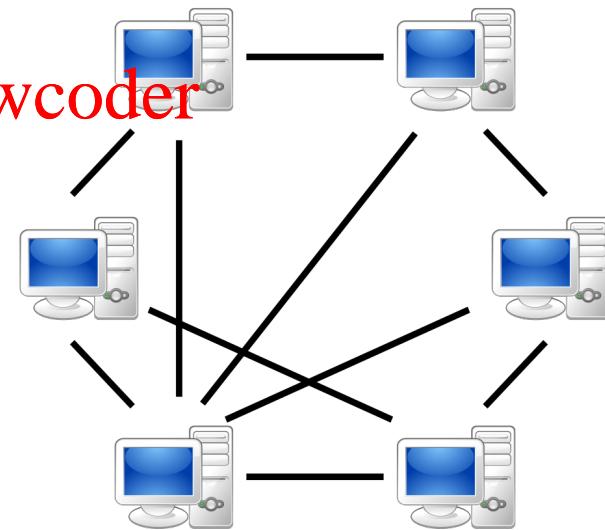


THE UNIVERSITY OF
SYDNEY

The peer-to-peer model

Every machine acts similarly

- Every machine is both a client and a server
- No centralized control: the responsibility is distributed evenly
Assignment Project Exam Help
<https://powcoder.com>
- Even the program executing on each machine is similar
- Forms an **overlay network**,
 - nodes are processors,
 - links possible communication channels.



Structured Peer-to-Peer Model

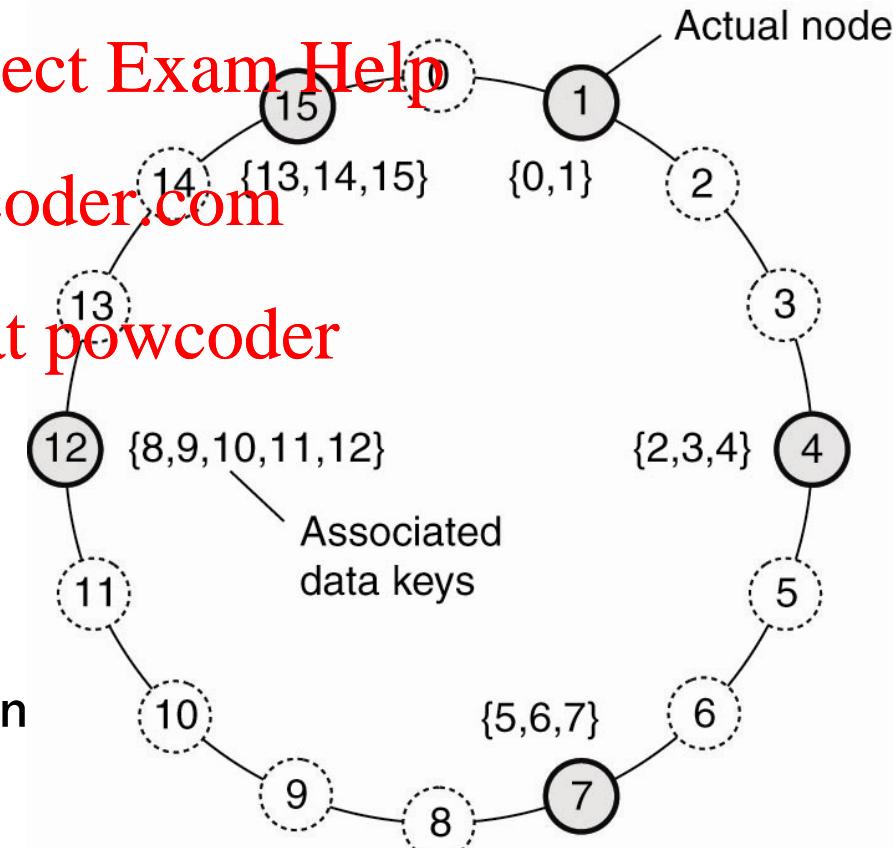
The overlay network is constructed using a deterministic procedure

- Chord is an example of a **Distributed Hash Table (DHT)**

As a node:

Assignment Project Exam Help

- I have a successor peer
<https://powcoder.com>
- I have a predecessor peer
- I have some **shortcuts** to other nodes to **speedup** delivery of requests
- I am responsible of a subset of the system data items (based on my unique identifier)



Hybrid Architectures

Assignment Project Exam Help

Architectures & Processes

Week 2, COMP3221

<https://powcoder.com>

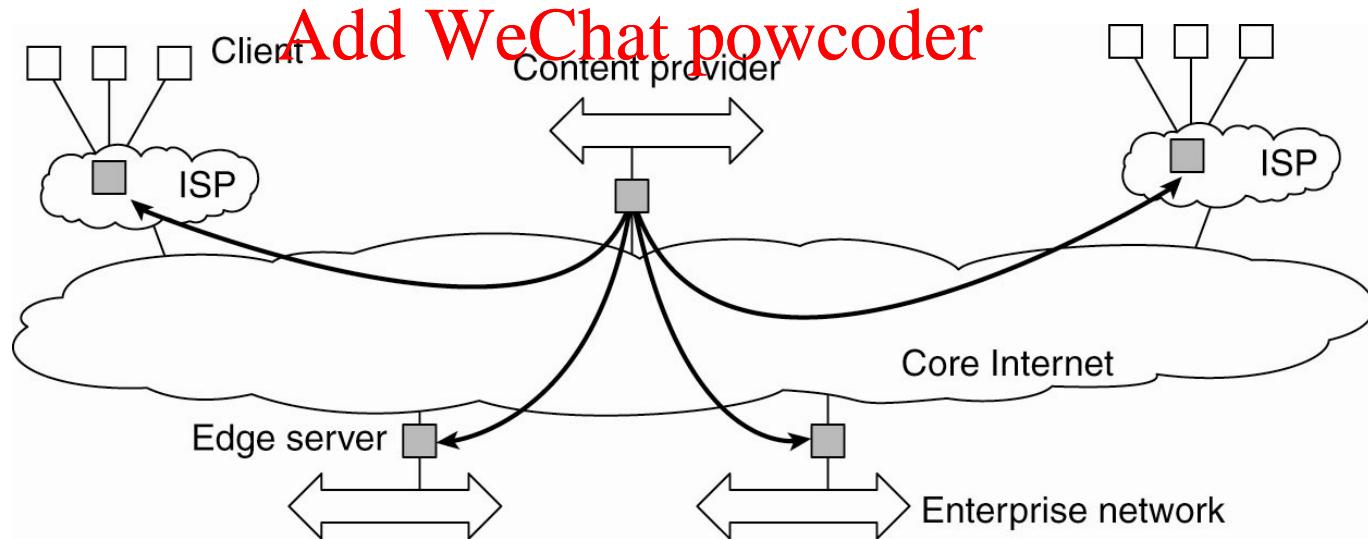
Add WeChat powcoder



THE UNIVERSITY OF
SYDNEY

Edge-Server Systems

- Servers are placed at the edge of the network.
- Edge servers serve content, possibly after applying filtering and transcoding functions
- One edge server acts as a origin server for all content, while taking advantage of other servers for optimized delivery of content, e.g. through replication.
<https://powcoder.com>



Collaborative Distributed Systems

- How do you first get started ?
 - Often a traditional client-server models is used.
 - Once joined fully decentralized schemes can be used.

Assignment Project Exam Help

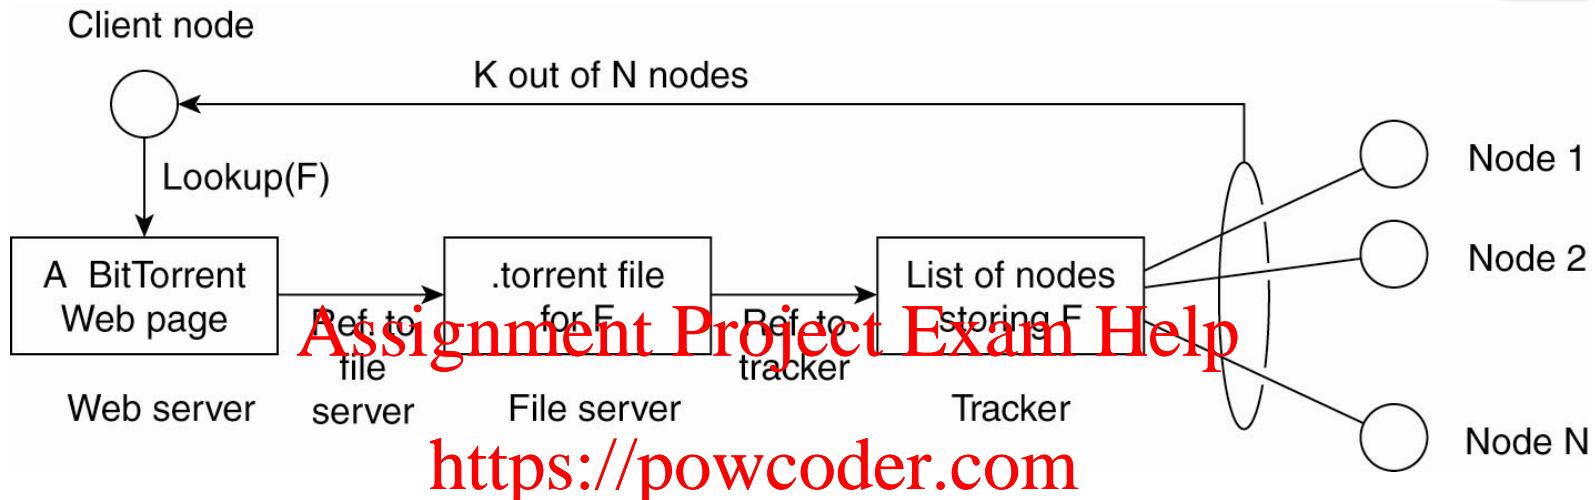
- Example: BitTorrent, a file sharing application

<https://powcoder.com>

- 35% of world-wide internet traffic in 2015.
- Used for Linux distribution, software patches, distributing movies
- Goal: quickly replicate large files to large number of clients



BitTorrent



- Web server hosts a .torrent file (w/ file length, hash, tracker's URL...)
- A tracker (server or a DHT) tracks downloaders/owners of a file
- Files are divided into chunks (256kb-1MB)
- Downloaders download chunks from themselves (and owners)
- **Tit-for-tat:** the more one shares (**server**), the faster it can download (**client**)

Conclusion

- Client and server are used to identify the role of communication participant
- Client and server roles may run on:
 - The same machine <https://powcoder.com>
 - Distinct machines with very different resources
 - Distinct machines with similar resources
- Most modern distributed systems are hybrid architectures

Processes

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Dr. Nguyen Tran
School of Computer Science



THE UNIVERSITY OF
SYDNEY

Introduction

- **Previously:** a distributed system gives the illusion of a single system to a user thanks to transparency

Assignment Project Exam Help

- **Now:** a single system gives the illusion of multiple resources to multiple users thanks to concurrency
<https://powcoder.com>

Add WeChat powcoder

Outline - Processes

- A Very Brief History
- UNIX Processes and Threads

Assignment Project Exam Help

- Multi-threading
<https://powcoder.com>
- Multi-threaded Server
Add WeChat powcoder

A Very Brief History

Architectures & Processes
Assignment Project Exam Help
Week 2, COMP3221
<https://powcoder.com>

Add WeChat powcoder



THE UNIVERSITY OF
SYDNEY

Evolution of Operating Systems (OSes)

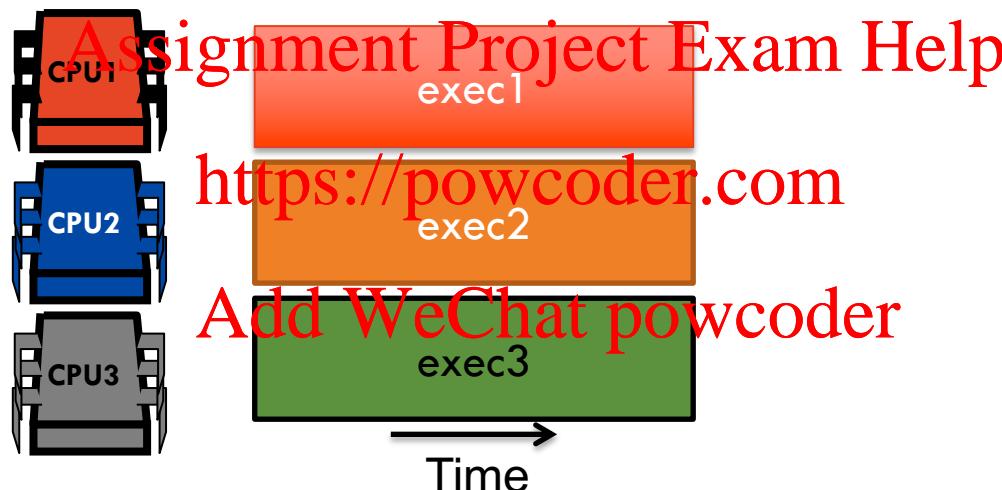
- The hardware evolved
 - Hardware was expensive, humans were cheap (e.g., Multics 1960's)
 - Hardware was cheaper, humans were expensive (e.g., desktop computer 1980's) **Assignment Project Exam Help**
 - Hardware are very cheap, humans are very expensive (e.g., handheld devices 2000's) <https://powcoder.com>
- The OS interaction had to adapt
 - Batch: one execution at a time
 - Multiprogramming: multiple program executions simultaneously
 - Timeshare: split time into slots allocated for different program execution
 - GUI: multiple interfaces to access a system from different end-points
 - Ubiquitous devices: each user possesses her own computational device

Uniprogramming vs. Multiprogramming

- **Uniprogramming:** one program execution at a time
 - MS/DOS
 - early Macintosh
 - Batch processing
 - No longer acceptable!
- **Multiprogramming:** more than one program execution at a time
 - Multics, UNIX/Linux, Windows NT, Mac OS X
- Many program executions on personal computers:
 - Browsing the web
 - Sending emails
 - Typing a letter
 - Burning a DVD

Concurrency

- Assume a single central processing unit (CPU)
- **Problem:** How to give the illusion to the users of multiple CPUs?

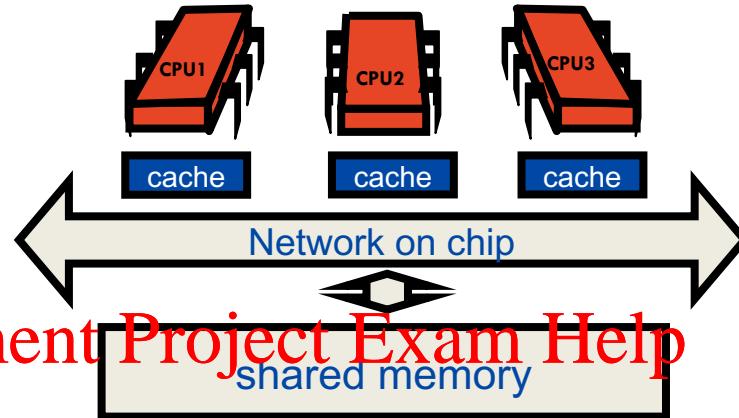


- **Solution:** **scheduling** program executions, one after the other during small fractions of the time



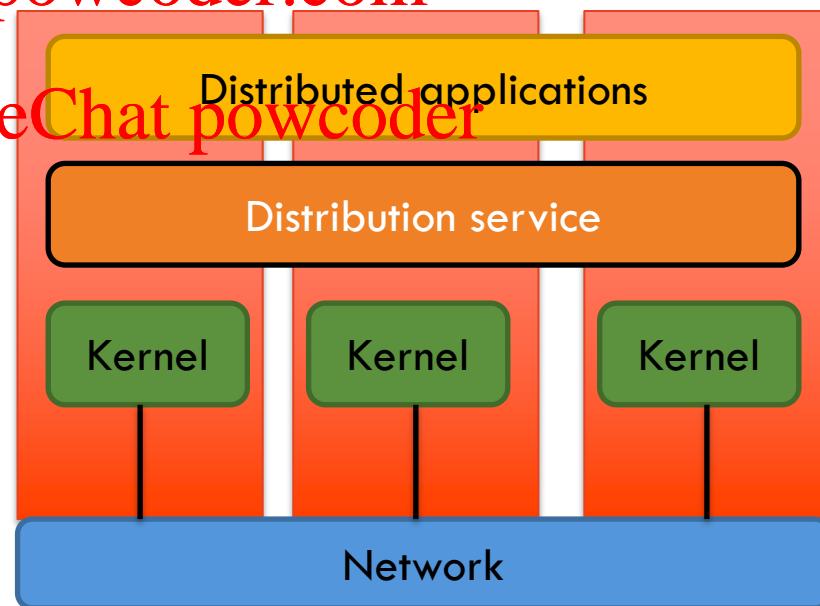
Concurrency

- Multi/many-cores on a single machine



<https://powcoder.com>

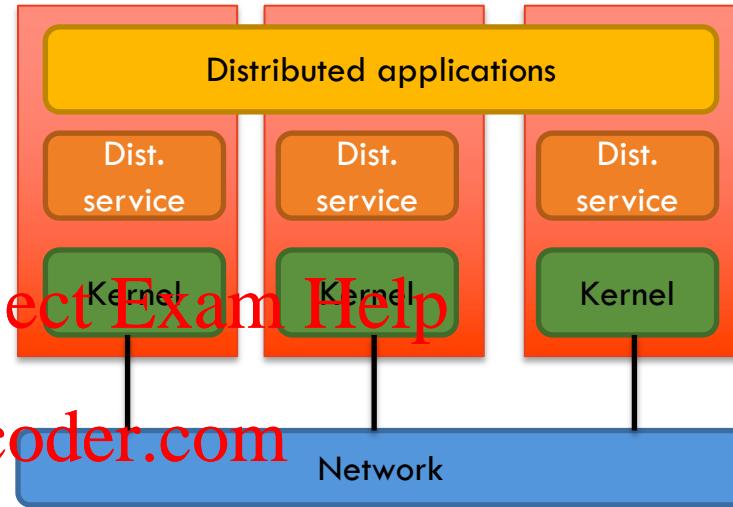
- Distributed operating systems
 - This is a single system image, the system maintains a single copy of the resources



Concurrency

- Network operating systems

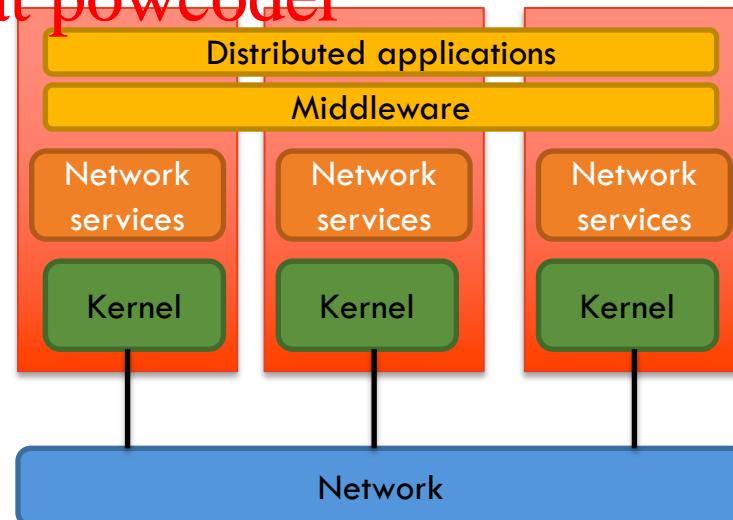
- Machines provide resources to other machines (e.g. UNIX rlogin)



- Middleware

- A layer over the network services providing general services to applications in a very transparent manner (systems can differ)

Add WeChat powcoder



Processes and Threads

Architectures & Processes
Assignment Project Exam Help
Week 2, COMP3221
<https://powcoder.com>

Add WeChat powcoder



THE UNIVERSITY OF
SYDNEY

Processes

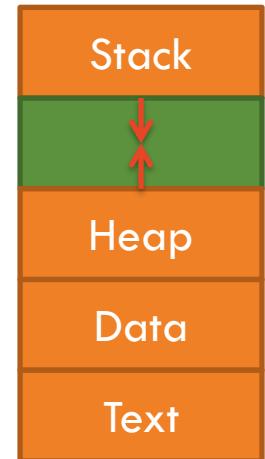
Process: defined as a program in execution

- Operating system creates a virtual CPU for each process
Assignment Project Exam Help
- OS makes sure that processes cannot maliciously or inadvertently affect each other's behaviour, i.e. **Secure !**
https://powcoder.com
- This concurrent transparency comes at a very high cost:
 - For every process (virtual CPU), OS must create a completely independent address space.

Address space

Address space: A unit of management of a process's virtual memory

- Process address space:
 - **Stack**: temporary data extensible towards lower virtual addresses
 - **Heap**: memory allocated dynamically extensible to higher virtual addresses
 - **Data** section: global variable
 - **Text** region: program code



Process

Process: an abstraction representing a program execution

- A process switches from a state to another:

- When there is timer interrupt goes off, explicit yield, I/O, etc.

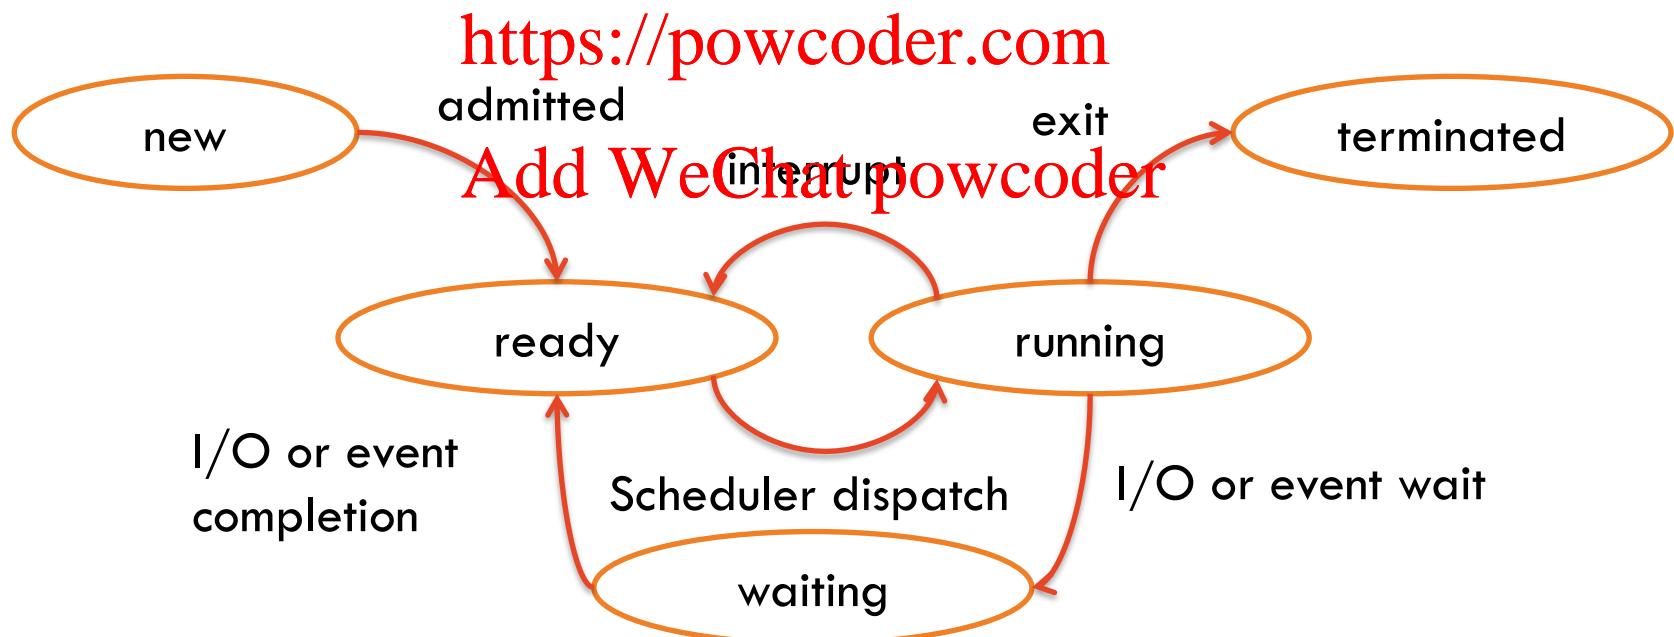
Assignment Project Exam Help

exec1

exec1.1

exec2.1

Time

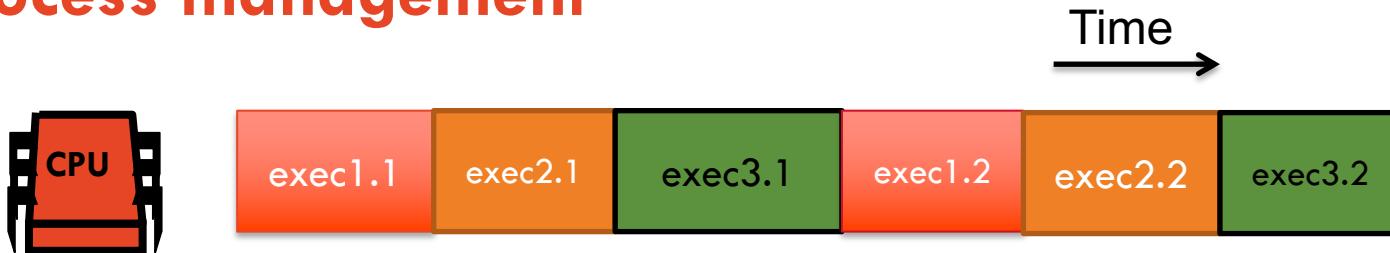


Process management

Process creation:

- Process allocation: the act of choosing the host of the process
 - The system provides command to execute a process at an idle workstation
 - The system chooses a host from a pool of processors to execute it
- Execution environment: an address space w/ initialized content & open files
 - Static: program text, heap and stack regions created from a list
 - Dynamic: UNIX fork shares program text and copies stack and heap regions

Process management



- Switching from one process to another has become one of the most expensive operations due to Context-Switches
 - CPU context: register values, program counter, stack pointer, etc.
 - Modify registers of MMU (Memory management unit)
 - Invalidate address translation caches such as in TLB (Translation lookaside buffer)

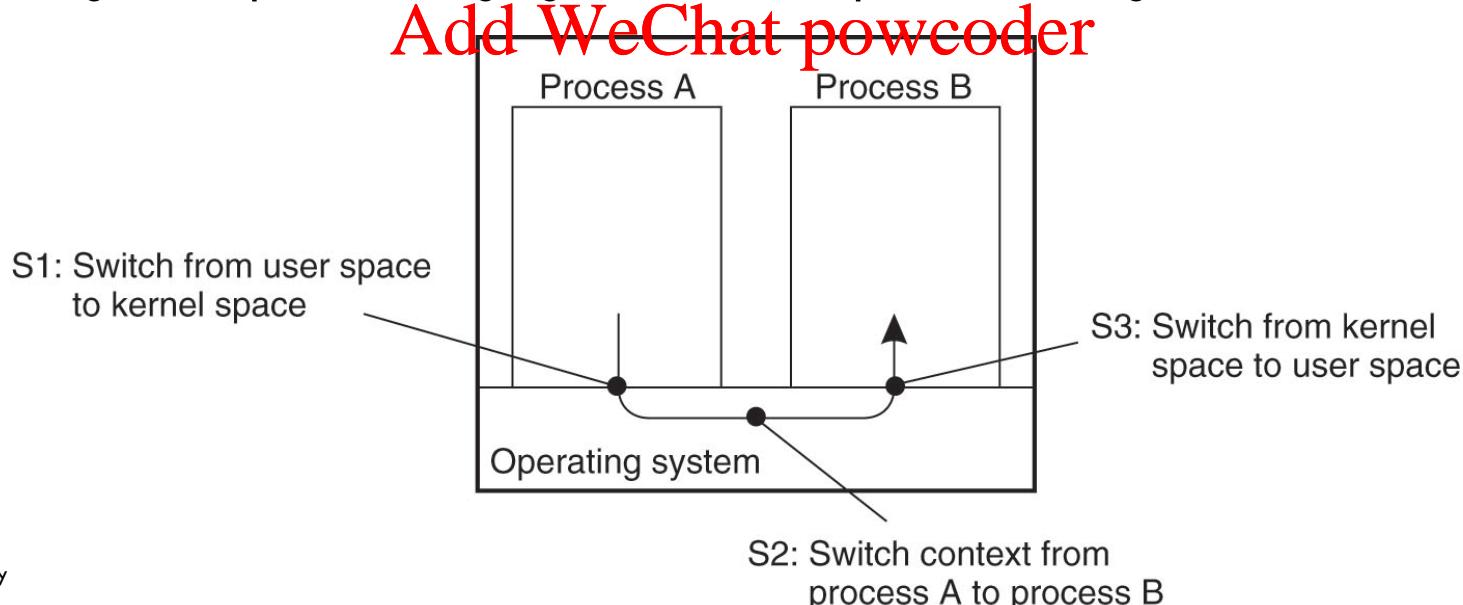
Inter-Process Communication (IPC):

- Pipes/Message queues/Shared memory segments
- Requires costly context switches

Context-switch

IPC requires kernel intervention

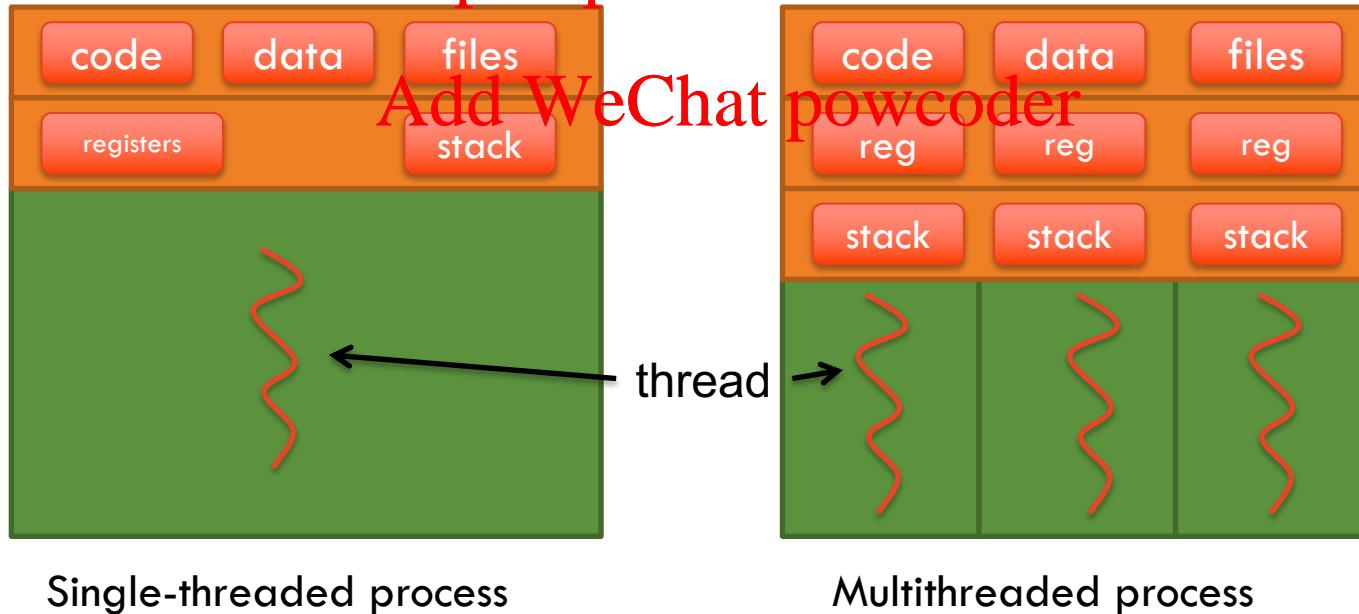
- S1. Process A switches from user to kernel mode
 - Changing the memory map in the MMU (memory management unit).
 - Flushing the TLB (translation lookaside buffer)
- S2. Process context-switch (swapping processes from A to B)
- S3. Process B switches from kernel to user mode
 - Again requires changing the MMU map and flushing the TLB



Thread

Threads: smallest unit of CPU utilization

- Multiple threads per process
- Portable OS Interface (POSIX) Threads
 - Current activity: program counter
 - Stack: temporary data
 - No data, code, files (but it shares them with other threads)
<https://powcoder.com>



Thread

- Portable OS Interface (POSIX) Threads
 - Current activity: program counter
 - Stack: temporary data
 - No data, code, files (but it shares them with other threads)
<https://powcoder.com>
- Communication between threads always through memory
 - Does not need IPC
 - No context switches required (when purely at user-level)

Thread Implementation

User-level threads library vs. kernel-scheduled thread

- **User-level thread library:**
 1. Cheap to create and destroy
 2. Blocking system call freezes the entire process of the thread (and other threads)
 3. Cheap context-switch:
 - Few instructions to store and reload CPU register values
 - No need to change memory map in MMU
 - No need to flush the TLB
- **Kernel-scheduled threads:**
 1. Costly to create and destroy
 2. Do not block the current process (and other threads) upon blocking system calls (I/O)
 3. Costly context-switch (similar to process context switch)
 - Needs to change memory map in MMU
 - Needs to flush the TLB

Processes vs. threads

- Processes
 - Isolated: prevents one process from interfering with another
 - Inefficient: starting/terminating a process and context switches are costly
- Threads
 - Non-isolated: avoiding incorrect interferences makes programming harder
 - Efficient: a thread is a lightweight version of a process
 - Multiple threads of control per process

<https://powcoder.com>

Add WeChat powcoder

Client multi-threading



Example: The Google Chrome web browser

- Multithreading: each tab runs its own thread
- A JavaScript error does not crash the main Chrome process:
Only one tab freezes, the user can close it independently of others
- Takes benefit of multicore architectures, each thread runs on a separate core
<https://powcoder.com>

A screenshot of the Google Chrome browser window. At the top, there is a toolbar with a 'New Tab' button and a search bar containing the URL 'http://www.google.com/we...'. Below the search bar is a bookmarks bar with the text 'For quick access, place your bookmarks here on the bookmarks bar. Import bookmarks now...'. In the center of the window, there are four tabs open: 'New Tab', 'http://www.google.com/we...', 'Assignment Project Exam Help', and 'Exam Help'. Below the tabs, there is a message 'Not signed in to Chrome (You're missing out—sign in)'. At the bottom of the window, there is a navigation bar with icons for 'Chrome ...', 'YouTube', 'Google Se...', and 'Gmail', along with links for 'Most visited' and 'Apps'. A large orange callout box on the right side of the window contains the text 'one distinct thread running each tab'.

Conclusion

- Concurrency has been used for decades since hardware was powerful enough to address multiple human needs

Assignment Project Exam Help

- Processes are heavy but protected whereas threads are lightweight but non-protected

Add WeChat powcoder

- Threads require synchronization to be protected from each other

What's Next ?

- Tutorial on Wednesday.
 - Multithreading
- Read Chapter 2 and Assignment Project Exam Help
- Next week: How do we communicate in distributed systems ?
<https://powcoder.com>
Add WeChat powcoder
- See you all next week !